


```
In [13]: #sentiment count
movie['sentiment'].value_counts()
```

```
Out[13]: positive    25000
         negative    25000
         Name: sentiment, dtype: int64
```

We can see that the dataset is balanced because it contains an equal number of positive and negative reviews.

```
In [14]: # plotting the labels using seaborn


sns.countplot(movie.sentiment)
plt.title('Countplot of the Sentiments', fontsize = 20, color = 'Violet')
plt.show()
```



```
In [59]: # Using seaborn to plot the sentence length
         # specifying the figure size
         # Labeling the title, x axis and y axis
         # Displaying the chart

plt.figure(figsize = (8,5), dpi = 100)
sns.distplot(lengths, color = 'brown')
plt.title('Histogram: Number of words in sentences', fontsize = 15, color = 'Blue')
plt.ylabel('Number of reviews')
plt.xlabel('Number of words')
plt.show()
```

Histogram: Number of words in sentences



```
In [69]: # Creating a simple deep neural network
# create an embedding layer by specifying the parameters we created earlier
# Add it to the model
# flatten the embedding layer since we are directly connecting it to densely connected layer
# At the end we add a dense layer with sigmoid activation function.

model = Sequential()

embedding_layer = Embedding(vocab_size, 100, weights = [embedding_matrix],
                           input_length = maxlen, trainable = False)

model.add(embedding_layer)

model.add(Flatten())

model.add(Dense(1, activation = 'sigmoid'))
```

```
In [70]: # Compiling our model
# Compile defines the loss function, the optimizer and the metrics.

model.compile(optimizer = 'adam', loss = 'binary_crossentropy',
             metrics = ['acc'])

# printing the summary of information about our model
print(model.summary())
```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 100, 100)	9254700
flatten_1 (Flatten)	(None, 10000)	0
dense_1 (Dense)	(None, 1)	10001

=====
 Total params: 9,264,701
 Trainable params: 10,001
 Non-trainable params: 9,254,700
 =====

None

```



In [77]: # Plotting the loss and accuracy differences for training and test sets
          # Using matplotlib plotting library
          # Giving a title to our chart
          # Labeling the x and y axis
          # creating a legend
          # displaying the chart

          plt.plot(history.history['acc'])
          plt.plot(history.history['val_acc'])

          plt.title('Model Accuracy')
          plt.ylabel('Accuracy')
          plt.xlabel('Epoch')
          plt.legend(['train', 'test'], loc = 'upper left')
          plt.show()

          plt.plot(history.history['loss'])
          plt.plot(history.history['val_loss'])

          plt.title('Model Loss')
          plt.ylabel('Loss')
          plt.xlabel('Epoch')
          plt.legend(['train', 'test'], loc = 'upper left')
          plt.show()
  
```



```
In [78]: # Building the model architecture
# Create an embedding layer by specifying the parameters we created earlier
# Creating a Recurrent neural network
# Here we will use LSTM (Long Term Short Term Memory)
# Bidirectional means the RNN processes sequence from start to end, and also backwards
# This makes the model perform better.
# We added another hidden layer and included an activation function as relu.s
# At the end we add a dense layer with sigmoid activation function.
#

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, 100,
                              weights=[embedding_matrix],
                              input_length=maxlength, trainable=False),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(128)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
In [79]: # Compiling the model
# Here we specify the loss, optimizer and metrics appropriately.

model.compile(loss='binary_crossentropy',
              optimizer=tf.keras.optimizers.Adam(1e-4),
              metrics=['accuracy'])
```

```
In [80]: # printing the model summary to view parameters
model.summary()
```

```
Model: "sequential"

Layer (type)                 Output Shape              Param #
-----
embedding (Embedding)        (None, 100, 100)         9254700
bidirectional (Bidirectional) (None, 256)              234496
dense (Dense)                 (None, 128)              32896
dense_1 (Dense)               (None, 1)                129
-----
Total params: 9,522,221
Trainable params: 267,521
Non-trainable params: 9,254,700
```

```
In [84]: # Making predictions

y_pred = model.predict_classes(X_test)

y_pred

# printing the classification report
print(classification_report(y_test, y_pred))
```

```
              precision    recall  f1-score   support

     0       0.79       0.84       0.81       4961
     1       0.83       0.78       0.80       5039

 accuracy          0.81
 macro avg         0.81       0.81       0.81
weighted avg         0.81       0.81       0.81
```


```
In [88]: # Plotting the loss and accuracy differences for training and test sets
# Using matplotlib plotting library
# Giving a title to our chart
# Labeling the x and y axis
# creating a legend
# displaying the chart
```


```
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])

plt.title('MODEL PERFORMANCE', fontsize = 20)
plt.ylabel('ACCURACY', fontsize = 15)
plt.xlabel('EPOCH', fontsize = 15)
plt.legend(['TRAIN DATA', 'TEST DATA'], loc = 'upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('Model Loss using RNN')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'test'], loc = 'upper left')
plt.show()
```





```
In [89]: # Building the model
# Create an embedding layer by specifying the parameters we created earlier
# Creating two Recurrent neural network layers
# Here we will use two LSTM (Long Term Short Term Memory) Layers
# Bidirectional means the RNN processes sequence from start to end, and also backwards
# This makes the model perform better.
# We added another hidden layer and included an activation function as relu.
# Here we added Dropout to prevent the model from over fitting
# Dropout randomly removes some neurons in the hidden layers
# At the end we add a dense layer with sigmoid activation function.

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, 100,
                              weights=[embedding_matrix],
                              input_length=maxlength, trainable=False),

    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(50, return_sequences = True)),

    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(25)),

    tf.keras.layers.Dense(50, activation='relu'),

    tf.keras.layers.Dropout(0.5),

    tf.keras.layers.Dense(1, activation='sigmoid')])
```

```
In [90]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 100, 100)	9254700
bidirectional_1 (Bidirection	(None, 100, 100)	60400
bidirectional_2 (Bidirection	(None, 50)	25200
dense_2 (Dense)	(None, 50)	2550
dropout (Dropout)	(None, 50)	0
dense_3 (Dense)	(None, 1)	51
Total params: 9,342,901		
Trainable params: 88,201		
Non-trainable params: 9,254,700		

```
In [103]: model = Sequential()
          embedding_layer = Embedding(vocab_size, 100, weights=[embedding_matrix], input_length=maxlength, trainable=False)
          model.add(embedding_layer)
          model.add(Conv1D(128, 5, activation='relu'))
          model.add(GlobalMaxPooling1D())
          model.add(Dropout(0.2)),
          model.add(Dense(1, activation='sigmoid'))
```

```
In [104]: model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['acc'])
```

```
In [105]: print(model.summary())
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 100, 100)	9254700
conv1d_1 (Conv1D)	(None, 96, 128)	64128
global_max_pooling1d_1 (Glob	(None, 128)	0
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 1)	129
Total params: 9,318,957		
Trainable params: 64,257		
Non-trainable params: 9,254,700		

None

Conclusion

- . Accuracy score on CNN is 83.21%
- . Accuracy score using Simple Deep Neural Networks is 72.35%
- . Accuracy using Recurrent Neural Network (RNN) with LSTM and Dropout is 80%
- . Accuracy using 2 RNN layer with LSTM is 82%
- . CNN model gives best accuracy among all with 83.21% on training and testing dataset.
- . Using the RNN (LSTM) with two hidden layers is also best which yield 82% accuracy on both training and testing dataset.
- . Neural Network models effective for sentiment analysis on IMDB reviews.

Optimizing the model further may yield better result using more data.